

ASDS-6303 Final Project

Default of Credit Card Clients

Phuong Trinh, Kyra Stolarski

2025-12-10

Load Packages

```
library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##     filter, lag

## The following objects are masked from 'package:base':
##     intersect, setdiff, setequal, union

library(ggplot2)
library(caret)

## Loading required package: lattice

library(pROC)

## Type 'citation("pROC")' for a citation.

##
## Attaching package: 'pROC'

## The following objects are masked from 'package:stats':
##     cov, smooth, var

library(rpart)
library(rpart.plot)
library(corrplot)

## corrplot 0.95 loaded
```

```

library(randomForest)

## randomForest 4.7-1.2

## Type rfNews() to see new features/changes/bug fixes.

## 
## Attaching package: 'randomForest'

## The following object is masked from 'package:ggplot2':
## 
##     margin

## The following object is masked from 'package:dplyr':
## 
##     combine

```

Load Dataset and set seed for reproducibility

```

set.seed(7)

credit <- read.csv("C:/Users/stolarskikm/Downloads/default+of+credit+card+clients/default of credit car
                     header = FALSE,
                     stringsAsFactors = FALSE)

head(credit)

##   V1      V2      V3      V4      V5      V6      V7      V8      V9      V10     V11     V12
## 1          X1      X2      X3      X4      X5      X6      X7      X8      X9      X10     X11
## 2 ID LIMIT_BAL SEX EDUCATION MARRIAGE AGE PAY_0 PAY_2 PAY_3 PAY_4 PAY_5 PAY_6
## 3 1    20000  2       2       1    24     2     2    -1    -1    -2    -2
## 4 2    120000 2       2       2    26    -1     2     0     0     0     0     2
## 5 3    90000  2       2       2    34     0     0     0     0     0     0     0
## 6 4    50000  2       2       1    37     0     0     0     0     0     0     0
##          V13      V14      V15      V16      V17      V18      V19      V20
## 1          X12      X13      X14      X15      X16      X17      X18      X19
## 2 BILL_AMT1 BILL_AMT2 BILL_AMT3 BILL_AMT4 BILL_AMT5 BILL_AMT6 PAY_AMT1 PAY_AMT2
## 3    3913     3102     689      0      0      0      0     689
## 4    2682     1725    2682    3272    3455    3261      0    1000
## 5   29239    14027   13559   14331   14948   15549    1518    1500
## 6   46990    48233   49291   28314   28959   29547    2000   2019
##          V21      V22      V23      V24          V25
## 1          X20      X21      X22      X23          Y
## 2 PAY_AMT3 PAY_AMT4 PAY_AMT5 PAY_AMT6 default payment next month
## 3      0      0      0      0          1
## 4    1000    1000      0    2000          1
## 5    1000    1000    1000    5000          0
## 6    1200    1100   1069    1000          0

```

Data Cleaning and Preprocessing

```
# Fix the header row
# Set first row as column names
colnames(credit) <- credit[2, ]

# Remove the first 2 rows
credit <- credit[-c(1,2), ]
head(credit)

##   ID LIMIT_BAL SEX EDUCATION MARRIAGE AGE PAY_0 PAY_2 PAY_3 PAY_4 PAY_5 PAY_6
## 3 1    20000  2      2        1  24     2     2    -1    -1    -2    -2
## 4 2    120000 2      2        2  26    -1     2     0     0     0     0
## 5 3    90000  2      2        2  34     0     0     0     0     0     0
## 6 4    50000  2      2        1  37     0     0     0     0     0     0
## 7 5    50000  1      2        1  57    -1     0    -1     0     0     0
## 8 6    50000  1      1        2  37     0     0     0     0     0     0
##   BILL_AMT1 BILL_AMT2 BILL_AMT3 BILL_AMT4 BILL_AMT5 BILL_AMT6 PAY_AMT1 PAY_AMT2
## 3      3913     3102      689       0       0       0       0      689
## 4      2682     1725     2682     3272     3455     3261       0     1000
## 5      29239    14027    13559    14331    14948    15549    1518     1500
## 6      46990    48233    49291    28314    28959    29547    2000     2019
## 7      8617     5670    35835    20940    19146    19131    2000    36681
## 8      64400    57069    57608    19394    19619    20024    2500     1815
##   PAY_AMT3 PAY_AMT4 PAY_AMT5 PAY_AMT6 default payment next month
## 3        0       0       0       0                   1
## 4      1000     1000      0     2000                   1
## 5      1000     1000     1000     5000                   0
## 6      1200     1100     1069     1000                   0
## 7     10000     9000      689      679                   0
## 8       657     1000     1000      800                   0

# Convert all character columns to numeric
credit <- credit %>% mutate(across(where(is.character), readr::parse_number))

# Sanity check
str(credit)

## 'data.frame': 30000 obs. of  25 variables:
## $ ID           : num  1 2 3 4 5 6 7 8 9 10 ...
## $ LIMIT_BAL    : num  20000 120000 90000 50000 50000 50000 500000 100000 140000 200000
## $ SEX          : num  2 2 2 2 1 1 1 2 2 1 ...
## $ EDUCATION    : num  2 2 2 2 2 1 1 2 3 3 ...
## $ MARRIAGE     : num  1 2 2 1 1 2 2 2 1 2 ...
## $ AGE          : num  24 26 34 37 57 37 29 23 28 35 ...
## $ PAY_0         : num  2 -1 0 0 -1 0 0 0 -2 ...
## $ PAY_2         : num  2 2 0 0 0 0 0 -1 0 -2 ...
## $ PAY_3         : num  -1 0 0 0 -1 0 0 -1 2 -2 ...
## $ PAY_4         : num  -1 0 0 0 0 0 0 0 -2 ...
## $ PAY_5         : num  -2 0 0 0 0 0 0 0 -1 ...
## $ PAY_6         : num  -2 2 0 0 0 0 0 -1 0 -1 ...
## $ BILL_AMT1    : num  3913 2682 29239 46990 8617 ...
```

```

## $ BILL_AMT2 : num 3102 1725 14027 48233 5670 ...
## $ BILL_AMT3 : num 689 2682 13559 49291 35835 ...
## $ BILL_AMT4 : num 0 3272 14331 28314 20940 ...
## $ BILL_AMT5 : num 0 3455 14948 28959 19146 ...
## $ BILL_AMT6 : num 0 3261 15549 29547 19131 ...
## $ PAY_AMT1 : num 0 0 1518 2000 2000 ...
## $ PAY_AMT2 : num 689 1000 1500 2019 36681 ...
## $ PAY_AMT3 : num 0 1000 1000 1200 10000 657 38000 0 432 0 ...
## $ PAY_AMT4 : num 0 1000 1000 1100 9000 ...
## $ PAY_AMT5 : num 0 0 1000 1069 689 ...
## $ PAY_AMT6 : num 0 2000 5000 1000 679 ...
## $ default payment next month: num 1 1 0 0 0 0 0 0 0 0 ...

```

```
head(credit)
```

```

## ID LIMIT_BAL SEX EDUCATION MARRIAGE AGE PAY_0 PAY_2 PAY_3 PAY_4 PAY_5 PAY_6
## 3 1 20000 2 2 1 24 2 2 -1 -1 -2 -2
## 4 2 120000 2 2 2 26 -1 2 0 0 0 0
## 5 3 90000 2 2 2 34 0 0 0 0 0 0
## 6 4 50000 2 2 1 37 0 0 0 0 0 0
## 7 5 50000 1 2 1 57 -1 0 -1 0 0 0
## 8 6 50000 1 1 2 37 0 0 0 0 0 0
## BILL_AMT1 BILL_AMT2 BILL_AMT3 BILL_AMT4 BILL_AMT5 BILL_AMT6 PAY_AMT1 PAY_AMT2
## 3 3913 3102 689 0 0 0 0 689
## 4 2682 1725 2682 3272 3455 3261 0 1000
## 5 29239 14027 13559 14331 14948 15549 1518 1500
## 6 46990 48233 49291 28314 28959 29547 2000 2019
## 7 8617 5670 35835 20940 19146 19131 2000 36681
## 8 64400 57069 57608 19394 19619 20024 2500 1815
## PAY_AMT3 PAY_AMT4 PAY_AMT5 PAY_AMT6 default payment next month
## 3 0 0 0 0 1
## 4 1000 1000 0 2000 1
## 5 1000 1000 1000 5000 0
## 6 1200 1100 1069 1000 0
## 7 10000 9000 689 679 0
## 8 657 1000 1000 800 0

```

```

# Removing ID Column
credit <- select(credit, -ID)

# Cleaning column names
names(credit) <- tolower(names(credit))
names(credit)[names(credit) == 'default payment next month'] <- 'default_next_month'
names(credit)[names(credit) == 'pay_0'] <- 'pay_1'
head(credit)

```

```

## limit_bal sex education marriage age pay_1 pay_2 pay_3 pay_4 pay_5 pay_6
## 3 20000 2 2 1 24 2 2 -1 -1 -2 -2
## 4 120000 2 2 2 26 -1 2 0 0 0 0
## 5 90000 2 2 2 34 0 0 0 0 0 0
## 6 50000 2 2 1 37 0 0 0 0 0 0
## 7 50000 1 2 1 57 -1 0 -1 0 0 0
## 8 50000 1 1 2 37 0 0 0 0 0 0

```

```

##   bill_amt1 bill_amt2 bill_amt3 bill_amt4 bill_amt5 bill_amt6 pay_amt1 pay_amt2
## 3      3913      3102       689        0        0        0        0       689
## 4      2682      1725      2682      3272      3455      3261        0     1000
## 5     29239     14027     13559     14331     14948     15549     1518     1500
## 6     46990     48233     49291     28314     28959     29547     2000     2019
## 7      8617      5670     35835     20940     19146     19131     2000   36681
## 8     64400     57069     57608     19394     19619     20024     2500     1815
##   pay_amt3 pay_amt4 pay_amt5 pay_amt6 default_next_month
## 3        0        0        0        0          1
## 4     1000     1000        0     2000          1
## 5     1000     1000     1000      5000          0
## 6     1200     1100     1069     1000          0
## 7    10000     9000       689       679          0
## 8      657     1000     1000       800          0

# Re-coding categorical attributes and target

# SEX: 1=Male, 2=Female
credit <- credit %>%
  mutate(sex = factor(sex, levels = c(1,2), labels = c("Male","Female")))

# EDUCATION: fix odd codes to "Others"
credit <- credit %>%
  mutate(education = dplyr::recode(education,
    `0`="Others", `1`="GradSchool", `2`="University",
    `3`="HighSchool", `4`="Others", `5`="Others", `6`="Others"),
  education = factor(education))

# MARRIAGE: 1=Married, 2=Single, others -> "Others"
credit <- credit %>%
  mutate(marriage = dplyr::recode(marriage,
    `0`="Others", `1`="Married", `2`="Single", `3`="Others"),
  marriage = factor(marriage))

# PAY_0..PAY_6 as ordered factors (payment status codes)
pay_cols <- c("pay_1","pay_2","pay_3","pay_4","pay_5","pay_6")
credit <- credit %>%
  mutate(across(all_of(pay_cols), ~factor(., ordered = TRUE)))

# Target Y to factor (0=No Default, 1=Default)
credit <- credit %>%
  mutate(default_next_month =
    factor(default_next_month, levels = c(0,1),
          labels = c("NoDefault","Default")))) %>%
  rename(y = default_next_month)

str(credit[, c("sex","education","marriage",pay_cols,"y")])

## 'data.frame': 30000 obs. of 10 variables:
## $ sex      : Factor w/ 2 levels "Male","Female": 2 2 2 2 1 1 1 2 2 1 ...
## $ education: Factor w/ 4 levels "GradSchool","HighSchool",...: 4 4 4 4 4 1 1 4 2 2 ...
## $ marriage : Factor w/ 3 levels "Married","Others",...: 1 3 3 1 1 3 3 3 1 3 ...
## $ pay_1    : Ord.factor w/ 11 levels "-2"<"-1"<"0"<...: 5 2 3 3 2 3 3 3 1 ...

```

```

## $ pay_2 : Ord.factor w/ 11 levels "-2"<"-1"<"0"<...: 5 5 3 3 3 3 3 2 3 1 ...
## $ pay_3 : Ord.factor w/ 11 levels "-2"<"-1"<"0"<...: 2 3 3 3 2 3 3 2 5 1 ...
## $ pay_4 : Ord.factor w/ 11 levels "-2"<"-1"<"0"<...: 2 3 3 3 3 3 3 3 3 1 ...
## $ pay_5 : Ord.factor w/ 10 levels "-2"<"-1"<"0"<...: 1 3 3 3 3 3 3 3 3 2 ...
## $ pay_6 : Ord.factor w/ 10 levels "-2"<"-1"<"0"<...: 1 4 3 3 3 3 3 2 3 2 ...
## $ y      : Factor w/ 2 levels "NoDefault","Default": 2 2 1 1 1 1 1 1 1 1 ...

```

Missing values per column

```

colSums(is.na(credit))

```

```

## limit_bal          sex education marriage       age     pay_1     pay_2     pay_3
##      0            0        0        0        0        0        0        0
##      pay_4      pay_5      pay_6 bill_amt1 bill_amt2 bill_amt3 bill_amt4 bill_amt5
##      0            0        0        0        0        0        0        0
## bill_amt6  pay_amt1  pay_amt2  pay_amt3  pay_amt4  pay_amt5  pay_amt6
##      0            0        0        0        0        0        0

```

Descriptive Statistics

```

summary(credit)

```

```

##   limit_bal          sex      education      marriage
## Min.    : 10000  Male   :11888  GradSchool:10585  Married:13659
## 1st Qu.: 50000  Female:18112  HighSchool: 4917  Others  : 377
## Median  : 140000                         Others   : 468  Single  :15964
## Mean    : 167484                         University:14030
## 3rd Qu.: 240000
## Max.   :1000000
##
##   age          pay_1      pay_2      pay_3
## Min.  :21.00  0 :14737  0 :15730  0 :15764
## 1st Qu.:28.00 -1 : 5686 -1 : 6050 -1 : 5938
## Median :34.00  1 : 3688  2 : 3927 -2 : 4085
## Mean   :35.49 -2 : 2759 -2 : 3782  2 : 3819
## 3rd Qu.:41.00  2 : 2667  3 : 326   3 : 240
## Max.   :79.00  3 : 322   4 : 99    4 : 76
##           (Other): 141  (Other): 86  (Other): 78
##
##   pay_4      pay_5      pay_6      bill_amt1
## 0 :16455  0 :16947  0 :16286  Min.  :-165580
## -1 : 5687 -1 : 5539 -1 : 5740  1st Qu.: 3559
## -2 : 4348 -2 : 4546 -2 : 4895  Median : 22382
## 2 : 3159  2 : 2626  2 : 2766  Mean   : 51223
## 3 : 180   3 : 178   3 : 184   3rd Qu.: 67091
## 4 : 69    4 : 84    4 : 49    Max.   : 964511
## (Other): 102 (Other): 80  (Other): 80
##
##   bill_amt2      bill_amt3      bill_amt4      bill_amt5
## Min.  :-69777  Min.  :-157264  Min.  :-170000  Min.  :-81334
## 1st Qu.: 2985  1st Qu.: 2666  1st Qu.: 2327  1st Qu.: 1763
## Median : 21200 Median : 20089 Median : 19052 Median : 18105
## Mean   : 49179 Mean   : 47013 Mean   : 43263 Mean   : 40311
## 3rd Qu.: 64006 3rd Qu.: 60165 3rd Qu.: 54506 3rd Qu.: 50191

```

```

##   Max.    :983931    Max.    :1664089    Max.    : 891586    Max.    :927171
##
##   bill_amt6          pay_amt1          pay_amt2          pay_amt3
##   Min.    :-339603    Min.    :     0    Min.    :     0    Min.    :     0
##   1st Qu.:    1256    1st Qu.:   1000    1st Qu.:    833    1st Qu.:   390
##   Median :   17071    Median :   2100    Median :   2009    Median :   1800
##   Mean   :  38872    Mean   :  5664    Mean   :  5921    Mean   :  5226
##   3rd Qu.:  49198    3rd Qu.:  5006    3rd Qu.:  5000    3rd Qu.:  4505
##   Max.   : 961664    Max.   :873552    Max.   :1684259    Max.   :896040
##
##   pay_amt4          pay_amt5          pay_amt6          y
##   Min.    :     0    Min.    :  0.0    Min.    :  0.0    NoDefault:23364
##   1st Qu.:   296    1st Qu.: 252.5    1st Qu.: 117.8    Default  : 6636
##   Median :  1500    Median : 1500.0    Median : 1500.0
##   Mean   :  4826    Mean   : 4799.4    Mean   : 5215.5
##   3rd Qu.:  4013    3rd Qu.: 4031.5    3rd Qu.: 4000.0
##   Max.   :621000    Max.   :426529.0    Max.   :528666.0
##

```

The descriptive analysis shows several important patterns in the credit card default dataset. The average credit limit (LIMIT_BAL) is about 167,000, but the distribution is highly skewed, with some clients having extremely high limits. Most customers fall within the 30–41 age range, with a median age of 34. The majority of clients are female, university educated, and single or married, reflecting the demographic structure of the dataset.

Payment history variables (PAY_0 to PAY_6) indicate that most customers are either on time or slightly delayed in previous months, though the presence of negative values (e.g., -1, -2) shows early payments or no usage. Bill amounts and payment amounts have strong right-skewness, with a few very large outliers across BILL_AMT and PAY_AMT variables.

Finally, the target variable shows that 23,335 clients (78%) are non-defaulters, while 6,630 clients (22%) defaulted, indicating moderate class imbalance. This imbalance is important to consider during modeling, especially for evaluation metrics and potential resampling techniques.

Exploratory Data Analysis

Default Distribution


```

library(ggplot2)

ggplot(credit, aes(x = y, fill = y)) +
  geom_bar() +
  scale_fill_manual(values = c("lightgreen", "lightpink")) +
  labs(title = "Distribution of Default vs No Default",
       x = "Default Status",
       y = "Count") +
  theme_minimal()

```


Distribution of Default vs No Default



Histogram of Credit Limit

```
ggplot(credit, aes(x = limit_bal)) +
  geom_histogram(bins = 30, fill = "skyblue", color = "white") +
  labs(title = "Distribution of Credit Limit",
       x = "Credit Limit (limit_bal)",
       y = "Frequency") +
  theme_minimal()
```


Distribution of Credit Limit



Credit Limit by Default Status


```
ggplot(credit, aes(x = y, y = limit_bal, fill = y)) +
  geom_boxplot() +
  scale_fill_manual(values = c("lightgreen", "lightpink")) +
  labs(title = "Credit Limit by Default Status",
       x = "Default Status",
       y = "Credit Limit") +
  theme_minimal()
```

Credit Limit by Default Status




Histogram of age

```
ggplot(credit, aes(x = age)) +  
  geom_histogram(bins = 30, fill = "darkorange", color = "white") +  
  labs(title = "Age Distribution",  
       x = "Age",  
       y = "Frequency") +  
  theme_minimal()
```



Age by Default

```
ggplot(credit, aes(x = y, y = age, fill = y)) +
  geom_boxplot() +
  scale_fill_manual(values = c("lightgreen", "lightpink")) +
  labs(title = "Age by Default Status",
       x = "Default Status",
       y = "Age") +
  theme_minimal()
```




Correlation Heatmap

```

num_vars <- dplyr::select(credit, dplyr::where(is.numeric))
corr_matrix <- stats::cor(num_vars, use = "pairwise.complete.obs")

corrplot::corrplot(
  corr_matrix,
  method = "color",
  type    = "lower",
  addCoef.col = "black",
  tl.col = "black",
  diag   = FALSE
)
  
```




```
# Find highly correlated variables & remove them
threshold <- 0.70
high_corr_vars <- caret::findCorrelation(
  corr_matrix,
  cutoff = threshold,
  names   = TRUE
)
print(high_corr_vars)
```

```
## [1] "bill_amt4" "bill_amt5" "bill_amt3" "bill_amt6" "bill_amt2"
```

```
credit <- credit[, !(names(credit) %in% high_corr_vars)]
```

Payment History vs Default

```
ggplot(credit, aes(x = pay_1, fill = y)) +
  geom_bar(position = "dodge") +
  labs(title = "PAY_1 Status vs Default",
       x = "Repayment Status (PAY_1)",
       y = "Count") +
  scale_fill_manual(values = c("lightgreen", "lightpink")) +
  theme_minimal()
```




PAY_0 = 0 or -1 → low default PAY_0 = 1 → strong default risk

This visually justifies why Decision Trees split on PAY_0 early

Payment Amount vs Bill Amount (Scatter)

```
ggplot(credit, aes(x = bill_amt1, y = pay_amt1)) +
  geom_point(alpha = 0.3, color = "red") +
  labs(title = "Bill Amount vs Payment Amount (Month 1)",
       x = "Bill Amount",
```

```
y = "Payment Amount") +
theme_minimal()
```



```
### Train-Test Split###
```

```
# Train-Test Split
train_index <- createDataPartition(credit$y, p = 0.7, list = FALSE)
```

```
train_set <- credit[train_index, ]
test_set <- credit[-train_index, ]
```

```
table(train_set$y)
```

```
##
## NoDefault    Default
##      16355      4646
```

```
# Oversampling Training Set
train_set_bal <- upSample(x = train_set[, -which(names(train_set) == "y")],
                           y = train_set$y,
                           yname = "y")
```

```
table(train_set_bal$y)
```

```
##
```

```
## NoDefault    Default
##      16355     16355
```

Models

```
# Function to plot confusion matrix
plot_confusion_matrix <- function(cfm, title, color_low, color_high) {
  cm_table <- as.data.frame(cfm$table)
  ggplot(cm_table, aes(x = Reference, y = Prediction)) +
    geom_tile(aes(fill = Freq), colour = "white") +
    geom_text(aes(label = Freq), size = 5) +
    scale_fill_gradient(low = color_low, high = color_high) +
    scale_y_discrete(limits = rev(levels(cm_table$Reference))) +
    ggtitle(title) +
    theme_minimal() +
    theme(
      plot.title = element_text(face = "bold", size = 14),
      axis.title = element_text(size = 12),
      axis.text = element_text(size = 11)
    )
}
```

Logistic Regression - Imbalanced Training

```
log_reg <- glm(y ~ ., family=binomial(link = 'logit'), data=train_set)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(log_reg)
```

```
##
## Call:
## glm(formula = y ~ ., family = binomial(link = "logit"), data = train_set)
##
## Coefficients: (1 not defined because of singularities)
##                  Estimate Std. Error z value Pr(>|z|)
## (Intercept)      3.369e+00  8.426e+01   0.040 0.968102
## limit_bal       -2.020e-06  2.095e-07  -9.644 < 2e-16 ***
## sexFemale        -1.416e-01  3.858e-02  -3.670 0.000242 ***
## educationHighSchool -6.460e-02  6.017e-02  -1.074 0.282981
## educationOthers   -8.521e-01  2.119e-01  -4.020 5.81e-05 ***
## educationUniversity 3.855e-04  4.466e-02   0.009 0.993112
## marriageOthers    -1.111e-01  1.660e-01  -0.669 0.503388
## marriageSingle     -1.408e-01  4.360e-02  -3.229 0.001241 **
## age                4.149e-03  2.369e-03   1.751 0.079861 .
## pay_1.L            4.237e+01  1.361e+03   0.031 0.975154
## pay_1.Q            4.022e+01  1.352e+03   0.030 0.976275
## pay_1.C            3.342e+01  1.064e+03   0.031 0.974932
## pay_1^4            2.283e+01  6.736e+02   0.034 0.972962
```

## pay_1^5	1.111e+01	3.297e+02	0.034	0.973129
## pay_1^6	3.037e+00	1.076e+02	0.028	0.977480
## pay_1^7	7.979e-01	2.215e+01	0.036	0.971270
## pay_1^8	-9.240e-01	2.911e+01	-0.032	0.974682
## pay_1^9	-5.177e-01	1.926e+01	-0.027	0.978552
## pay_1^10	1.623e-01	6.893e+00	0.024	0.981219
## pay_2.L	-3.075e+01	1.121e+03	-0.027	0.978126
## pay_2.Q	-1.037e+01	5.846e+02	-0.018	0.985841
## pay_2.C	1.881e+01	4.927e+02	0.038	0.969546
## pay_2^4	4.100e+01	1.065e+03	0.039	0.969278
## pay_2^5	4.658e+01	1.207e+03	0.039	0.969207
## pay_2^6	3.823e+01	9.473e+02	0.040	0.967811
## pay_2^7	2.354e+01	5.419e+02	0.043	0.965346
## pay_2^8	1.062e+01	2.161e+02	0.049	0.960807
## pay_2^9	2.534e+00	4.890e+01	0.052	0.958675
## pay_2^10	NA	NA	NA	NA
## pay_3.L	1.787e+01	3.162e+02	0.057	0.954932
## pay_3.Q	1.386e+01	3.294e+02	0.042	0.966436
## pay_3.C	2.097e+00	3.475e+02	0.006	0.995185
## pay_3^4	4.597e+00	2.378e+02	0.019	0.984576
## pay_3^5	8.727e+00	2.260e+02	0.039	0.969201
## pay_3^6	3.233e+00	2.145e+02	0.015	0.987974
## pay_3^7	8.184e+00	1.364e+02	0.060	0.952147
## pay_3^8	1.195e+01	2.521e+02	0.047	0.962177
## pay_3^9	-1.897e+00	2.711e+02	-0.007	0.994416
## pay_3^10	4.749e+00	1.522e+02	0.031	0.975103
## pay_4.L	-5.403e+01	5.227e+02	-0.103	0.917676
## pay_4.Q	-4.068e+01	4.512e+02	-0.090	0.928166
## pay_4.C	-1.066e+01	3.660e+02	-0.029	0.976759
## pay_4^4	-3.074e-01	2.397e+02	-0.001	0.998977
## pay_4^5	-6.367e-01	2.510e+02	-0.003	0.997976
## pay_4^6	3.819e+00	2.469e+02	0.015	0.987660
## pay_4^7	-5.172e+00	1.773e+02	-0.029	0.976725
## pay_4^8	-1.212e+01	2.673e+02	-0.045	0.963832
## pay_4^9	2.157e+00	2.762e+02	0.008	0.993768
## pay_4^10	-5.278e+00	1.535e+02	-0.034	0.972572
## pay_5.L	5.371e+01	8.507e+02	0.063	0.949657
## pay_5.Q	3.571e+01	7.437e+02	0.048	0.961705
## pay_5.C	9.128e+00	5.413e+02	0.017	0.986544
## pay_5^4	-8.506e+00	3.862e+02	-0.022	0.982426
## pay_5^5	-1.105e+01	2.575e+02	-0.043	0.965780
## pay_5^6	-4.037e+00	1.714e+02	-0.024	0.981207
## pay_5^7	2.206e+00	1.462e+02	0.015	0.987965
## pay_5^8	3.520e+00	1.071e+02	0.033	0.973790
## pay_5^9	2.040e+00	4.850e+01	0.042	0.966446
## pay_6.L	-7.630e+00	3.990e+02	-0.019	0.984743
## pay_6.Q	-5.594e+00	4.023e+02	-0.014	0.988905
## pay_6.C	-1.034e+00	3.332e+02	-0.003	0.997523
## pay_6^4	3.069e+00	2.437e+02	0.013	0.989951
## pay_6^5	4.073e+00	1.655e+02	0.025	0.980365
## pay_6^6	3.801e+00	1.049e+02	0.036	0.971092
## pay_6^7	3.259e+00	5.857e+01	0.056	0.955622
## pay_6^8	1.613e+00	2.620e+01	0.062	0.950927
## pay_6^9	3.559e-01	7.977e+00	0.045	0.964410

```

## bill_amt1          2.331e-06  3.691e-07   6.315 2.70e-10 ***
## pay_amt1           -1.322e-05 2.935e-06  -4.504 6.66e-06 ***
## pay_amt2           -9.339e-06 2.507e-06  -3.725 0.000195 ***
## pay_amt3           -3.316e-06 1.918e-06  -1.729 0.083892 .
## pay_amt4           -1.742e-06 1.856e-06  -0.938 0.347990
## pay_amt5           -2.605e-06 1.725e-06  -1.510 0.131053
## pay_amt6           -2.534e-06 1.588e-06  -1.595 0.110624
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 22196  on 21000  degrees of freedom
## Residual deviance: 18253  on 20928  degrees of freedom
## AIC: 18399
##
## Number of Fisher Scoring iterations: 12

# Training performance
log_pred_train <- predict(log_reg, newdata = train_set, type = "response")
predicted_train <- factor(ifelse(log_pred_train > 0.5, 1, 0), levels = c(0,1), labels = c("NoDefault","Default"))

con_mat <- confusionMatrix(predicted_train, train_set$y, mode = 'everything', positive = 'Default')
print(con_mat)

## Confusion Matrix and Statistics
##
##             Reference
## Prediction  NoDefault Default
##     NoDefault      15598    2994
##     Default         757     1652
##
##             Accuracy : 0.8214
##                 95% CI : (0.8161, 0.8265)
##     No Information Rate : 0.7788
##     P-Value [Acc > NIR] : < 2.2e-16
##
##             Kappa : 0.3737
##
## McNemar's Test P-Value : < 2.2e-16
##
##             Sensitivity : 0.35557
##             Specificity  : 0.95371
##     Pos Pred Value : 0.68576
##     Neg Pred Value : 0.83896
##             Precision  : 0.68576
##             Recall    : 0.35557
##             F1        : 0.46832
##             Prevalence : 0.22123
##     Detection Rate : 0.07866
##     Detection Prevalence : 0.11471
##     Balanced Accuracy : 0.65464
##
##     'Positive' Class : Default

```

```

##  

# Test prediction  

log_pred_test <- predict(log_reg, newdata = test_set, type = "response")  

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type == :  

## prediction from rank-deficient fit; attr(*, "non-estim") has doubtful cases  

predicted_test <- factor(ifelse(log_pred_test > 0.5, 1, 0), levels = c(0,1), labels = c("NoDefault","De  

con_mat <- confusionMatrix(predicted_test, test_set$y, mode = 'everything', positive = 'Default')  

print(con_mat)  

## Confusion Matrix and Statistics  

##  

##           Reference  

## Prediction  NoDefault Default  

##   NoDefault      6679     1265  

##   Default        330      725  

##  

##           Accuracy : 0.8228  

##                 95% CI : (0.8147, 0.8306)  

##   No Information Rate : 0.7789  

##   P-Value [Acc > NIR] : < 2.2e-16  

##  

##           Kappa : 0.3814  

##  

## McNemar's Test P-Value : < 2.2e-16  

##  

##           Sensitivity : 0.36432  

##           Specificity : 0.95292  

##   Pos Pred Value : 0.68720  

##   Neg Pred Value : 0.84076  

##           Precision : 0.68720  

##           Recall : 0.36432  

##           F1 : 0.47619  

##           Prevalence : 0.22114  

##   Detection Rate : 0.08056  

## Detection Prevalence : 0.11724  

##   Balanced Accuracy : 0.65862  


##  

## 'Positive' Class : Default  

##

```

Imbalanced Logistic Regression – Testing



```
## Setting levels: control = NoDefault, case = Default
```

```
## Setting direction: controls < cases
```

Logistic Regression - Balanced Training

```
log_reg_bal <- glm(y ~ ., family=binomial(link = 'logit'), data=train_set_bal)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(log_reg_bal)
```

```
##  
## Call:  
## glm(formula = y ~ ., family = binomial(link = "logit"), data = train_set_bal)  
##  
## Coefficients: (1 not defined because of singularities)  
##                                     Estimate Std. Error z value Pr(>|z|)  
## (Intercept)        4.552e+00  1.028e+02   0.044  0.964675  
## limit_bal       -2.017e-06  1.338e-07 -15.080  < 2e-16 ***  
## sexFemale      -1.159e-01  2.587e-02  -4.482 7.40e-06 ***  
## educationHighSchool -9.697e-03  4.011e-02  -0.242 0.808951
```

```

## educationOthers      -6.234e-01  1.235e-01 -5.049 4.45e-07 ***
## educationUniversity  3.954e-02  2.969e-02  1.332 0.182923
## marriageOthers       -9.179e-02  1.107e-01 -0.829 0.406902
## marriageSingle       -1.502e-01  2.905e-02 -5.169 2.35e-07 ***
## age                  2.947e-03  1.588e-03  1.856 0.063507 .
## pay_1.L               5.021e+01  2.231e+03  0.023 0.982041
## pay_1.Q               4.843e+01  2.217e+03  0.022 0.982572
## pay_1.C               4.030e+01  1.743e+03  0.023 0.981550
## pay_1^4                2.752e+01  1.103e+03  0.025 0.980088
## pay_1^5                1.371e+01  5.385e+02  0.025 0.979686
## pay_1^6                4.312e+00  1.737e+02  0.025 0.980194
## pay_1^7                1.210e+00  3.260e+01  0.037 0.970391
## pay_1^8                -8.603e-01  4.729e+01 -0.018 0.985488
## pay_1^9                -6.908e-01  3.152e+01 -0.022 0.982514
## pay_1^10               1.376e-01  1.130e+01  0.012 0.990280
## pay_2.L               -3.708e+01  1.832e+03 -0.020 0.983856
## pay_2.Q               -1.352e+01  9.423e+02 -0.014 0.988557
## pay_2.C               2.108e+01  7.942e+02  0.027 0.978828
## pay_2^4                4.753e+01  1.744e+03  0.027 0.978255
## pay_2^5                5.418e+01  1.979e+03  0.027 0.978160
## pay_2^6                4.443e+01  1.554e+03  0.029 0.977197
## pay_2^7                2.730e+01  8.894e+02  0.031 0.975508
## pay_2^8                1.229e+01  3.547e+02  0.035 0.972358
## pay_2^9                3.231e+00  8.024e+01  0.040 0.967874
## pay_2^10               NA          NA          NA          NA
## pay_3.L               1.913e+01  2.905e+02  0.066 0.947484
## pay_3.Q               1.485e+01  2.876e+02  0.052 0.958824
## pay_3.C               1.245e+00  3.914e+02  0.003 0.997461
## pay_3^4                4.654e+00  2.158e+02  0.022 0.982790
## pay_3^5                9.688e+00  3.022e+02  0.032 0.974426
## pay_3^6                2.974e+00  3.240e+02  0.009 0.992675
## pay_3^7                8.714e+00  1.920e+02  0.045 0.963807
## pay_3^8                1.327e+01  4.033e+02  0.033 0.973761
## pay_3^9                -2.998e+00  4.416e+02 -0.007 0.994585
## pay_3^10               5.505e+00  2.490e+02  0.022 0.982360
## pay_4.L               -5.867e+01  6.160e+02 -0.095 0.924113
## pay_4.Q               -4.371e+01  5.706e+02 -0.077 0.938935
## pay_4.C               -1.092e+01  5.437e+02 -0.020 0.983977
## pay_4^4                1.440e-01  3.613e+02  0.000 0.999682
## pay_4^5                -6.530e-01  3.725e+02 -0.002 0.998601
## pay_4^6                3.759e+00  3.575e+02  0.011 0.991611
## pay_4^7                -6.086e+00  2.236e+02 -0.027 0.978284
## pay_4^8                -1.345e+01  4.131e+02 -0.033 0.974022
## pay_4^9                2.269e+00  4.450e+02  0.005 0.995931
## pay_4^10               -5.550e+00  2.499e+02 -0.022 0.982281
## pay_5.L               5.690e+01  8.548e+02  0.067 0.946927
## pay_5.Q               3.742e+01  7.884e+02  0.047 0.962143
## pay_5.C               9.326e+00  6.172e+02  0.015 0.987945
## pay_5^4                -8.866e+00  4.536e+02 -0.020 0.984405
## pay_5^5                -1.125e+01  2.966e+02 -0.038 0.969743
## pay_5^6                -3.928e+00  1.801e+02 -0.022 0.982605
## pay_5^7                2.486e+00  1.363e+02  0.018 0.985452
## pay_5^8                3.530e+00  9.723e+01  0.036 0.971037
## pay_5^9                1.852e+00  4.396e+01  0.042 0.966402

```

```

## pay_6.L      -7.334e+00  4.274e+02 -0.017  0.986309
## pay_6.Q      -5.281e+00  4.392e+02 -0.012  0.990405
## pay_6.C      -9.789e-01  3.715e+02 -0.003  0.997898
## pay_6^4       2.876e+00  2.736e+02  0.011  0.991614
## pay_6^5       3.950e+00  1.807e+02  0.022  0.982555
## pay_6^6       3.666e+00  1.077e+02  0.034  0.972831
## pay_6^7       3.142e+00  5.618e+01  0.056  0.955408
## pay_6^8       1.580e+00  2.383e+01  0.066  0.947133
## pay_6^9       3.934e-01  6.997e+00  0.056  0.955164
## bill_amt1    2.103e-06  2.374e-07  8.858  < 2e-16 ***
## pay_amt1     -9.801e-06  1.486e-06 -6.594  4.28e-11 ***
## pay_amt2     -7.240e-06  1.309e-06 -5.531  3.18e-08 ***
## pay_amt3     -3.713e-06  1.036e-06 -3.584  0.000338 ***
## pay_amt4     -2.332e-06  1.140e-06 -2.045  0.040817 *
## pay_amt5     -2.896e-06  1.088e-06 -2.662  0.007777 **
## pay_amt6     -2.670e-06  9.920e-07 -2.691  0.007116 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 45346  on 32709  degrees of freedom
## Residual deviance: 37291  on 32637  degrees of freedom
## AIC: 37437
##
## Number of Fisher Scoring iterations: 13

```

```

# Balanced Training performance on original training set
log_pred_train_bal <- predict(log_reg_bal, newdata = train_set, type = "response")

```

```

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type == :
## prediction from rank-deficient fit; attr(*, "non-estim") has doubtful cases

predicted_train_bal <- factor(ifelse(log_pred_train_bal > 0.5, 1, 0), levels = c(0,1), labels = c("NoDe
con_mat <- confusionMatrix(predicted_train_bal, train_set$y, mode = 'everything', positive = 'Default')
print(con_mat)

```

```

## Confusion Matrix and Statistics
##
##             Reference
## Prediction  NoDefault Default
##   NoDefault      13708     1993
##   Default        2647     2653
##
##             Accuracy : 0.7791
##                 95% CI : (0.7734, 0.7847)
##   No Information Rate : 0.7788
##   P-Value [Acc > NIR] : 0.4642
##
##             Kappa : 0.3896
##
##   Mcnemar's Test P-Value : <2e-16

```

```

##          Sensitivity : 0.5710
##          Specificity : 0.8382
##          Pos Pred Value : 0.5006
##          Neg Pred Value : 0.8731
##          Precision : 0.5006
##          Recall : 0.5710
##          F1 : 0.5335
##          Prevalence : 0.2212
##          Detection Rate : 0.1263
##          Detection Prevalence : 0.2524
##          Balanced Accuracy : 0.7046
##
##          'Positive' Class : Default
##         

# Balanced testing performance
log_pred_test_bal <- predict(log_reg_bal, newdata = test_set, type = "response")

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type == :
## prediction from rank-deficient fit; attr(*, "non-estim") has doubtful cases

predicted_test_bal <- factor(ifelse(log_pred_test_bal > 0.5, 1, 0), levels = c(0,1), labels = c("NoDefa


con_mat <- confusionMatrix(predicted_test_bal, test_set$y, mode = 'everything', positive = 'Default')
print(con_mat)

## Confusion Matrix and Statistics
##
##          Reference
## Prediction  NoDefault Default
##  NoDefault      5873     830
##  Default        1136    1160
##
##          Accuracy : 0.7815
##          95% CI : (0.7728, 0.79)
##          No Information Rate : 0.7789
##          P-Value [Acc > NIR] : 0.2758
##
##          Kappa : 0.3989
##
##          Mcnemar's Test P-Value : 6.039e-12
##
##          Sensitivity : 0.5829
##          Specificity : 0.8379
##          Pos Pred Value : 0.5052
##          Neg Pred Value : 0.8762
##          Precision : 0.5052
##          Recall : 0.5829
##          F1 : 0.5413
##          Prevalence : 0.2211
##          Detection Rate : 0.1289
##          Detection Prevalence : 0.2551

```

```
##      Balanced Accuracy : 0.7104
##      'Positive' Class : Default
##
```


Balanced Logistic Regression – Testing



```
## Setting levels: control = NoDefault, case = Default
## Setting direction: controls < cases
```

CART Decision tree - Imbalanced Training Set


Imbalanced Decision Tree Model



```
print(dt_model)
```

```
## n= 21001
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
## 1) root 21001 4646 NoDefault (0.7787724 0.2212276)
##    2) pay_1=-2,-1,0,1 18827 3134 NoDefault (0.8335369 0.1664631) *
##    3) pay_1=2,3,4,5,6,7,8 2174  662 Default (0.3045078 0.6954922) *
```

Imbalanced Variable Importance



```
# Training performance
dt_pred_train <- predict(dt_model, train_set, type = "class")
con_mat <- confusionMatrix(dt_pred_train, train_set$y, mode = 'everything', positive = "Default")
print(con_mat)

## Confusion Matrix and Statistics
##
##             Reference
## Prediction  NoDefault Default
##   NoDefault      15693     3134
##   Default         662      1512
##
##                   Accuracy : 0.8192
##                   95% CI : (0.814, 0.8244)
##   No Information Rate : 0.7788
##   P-Value [Acc > NIR] : < 2.2e-16
##
##                   Kappa : 0.352
##
##   Mcnemar's Test P-Value : < 2.2e-16
##
##                   Sensitivity : 0.3254
##                   Specificity  : 0.9595
##   Pos Pred Value  : 0.6955
##   Neg Pred Value : 0.8335
```

```


##                  Precision : 0.6955
##                  Recall   : 0.3254
##                  F1      : 0.4434
##      Prevalence   : 0.2212
##      Detection Rate : 0.0720
##      Detection Prevalence : 0.1035
##      Balanced Accuracy : 0.6425
##
##      'Positive' Class : Default
##

# Testing performance
dt_pred_test <- predict(dt_model, test_set, type = "class")
con_mat <- confusionMatrix(dt_pred_test, test_set$y, mode = 'everything', positive = "Default")
print(con_mat)

## Confusion Matrix and Statistics
##
##      Reference
## Prediction  NoDefault Default
##    NoDefault     6718     1325
##    Default       291      665
##
##                  Accuracy : 0.8204
##                  95% CI  : (0.8123, 0.8283)
##      No Information Rate : 0.7789
##      P-Value [Acc > NIR] : < 2.2e-16
##
##                  Kappa : 0.3595
##
##      Mcnemar's Test P-Value : < 2.2e-16
##
##      Sensitivity : 0.3342
##      Specificity  : 0.9585
##      Pos Pred Value : 0.6956
##      Neg Pred Value : 0.8353
##                  Precision : 0.6956
##                  Recall   : 0.3342
##                  F1      : 0.4515
##      Prevalence   : 0.2211
##      Detection Rate : 0.0739
##      Detection Prevalence : 0.1062
##      Balanced Accuracy : 0.6463
##
##      'Positive' Class : Default
##

```

Imbalanced Decision Tree – Testing




```
## Setting levels: control = NoDefault, case = Default
```


```
## Setting direction: controls < cases
```

CART Decision tree - Balanced Training Set

Balanced Decision Tree Model



Balanced Variable Importance



```
# Balanced Training performance on original training set
dt_pred_train_bal <- predict(dt_model_bal, train_set, type = "class")
con_mat <- confusionMatrix(dt_pred_train_bal, train_set$y, mode = 'everything', positive = "Default")
print(con_mat)

## Confusion Matrix and Statistics
##
##             Reference
## Prediction  NoDefault Default
##   NoDefault      13972    2270
##   Default        2383     2376
##
##                   Accuracy : 0.7784
##                   95% CI : (0.7728, 0.784)
##   No Information Rate : 0.7788
##   P-Value [Acc > NIR] : 0.5502
##
##                   Kappa : 0.3625
##
##   Mcnemar's Test P-Value : 0.1006
##
##                   Sensitivity : 0.5114
##                   Specificity  : 0.8543
##   Pos Pred Value  : 0.4993
##   Neg Pred Value : 0.8602
```

```


##                  Precision : 0.4993
##                  Recall   : 0.5114
##                  F1      : 0.5053
##      Prevalence   : 0.2212
##      Detection Rate : 0.1131
##      Detection Prevalence : 0.2266
##      Balanced Accuracy : 0.6829
##
##      'Positive' Class : Default
##

# Testing performance
dt_pred_test <- predict(dt_model_bal, test_set, type = "class")
con_mat <- confusionMatrix(dt_pred_test, test_set$y, mode = 'everything', positive = "Default")
print(con_mat)

## Confusion Matrix and Statistics
##
##      Reference
## Prediction  NoDefault Default
##    NoDefault      6003     937
##    Default        1006    1053
##
##                  Accuracy : 0.7841
##                  95% CI  : (0.7754, 0.7926)
##      No Information Rate : 0.7789
##      P-Value [Acc > NIR] : 0.1186
##
##                  Kappa : 0.3809
##
##      Mcnemar's Test P-Value : 0.1229
##
##      Sensitivity : 0.5291
##      Specificity : 0.8565
##      Pos Pred Value : 0.5114
##      Neg Pred Value : 0.8650
##                  Precision : 0.5114
##                  Recall   : 0.5291
##                  F1      : 0.5201
##      Prevalence   : 0.2211
##      Detection Rate : 0.1170
##      Detection Prevalence : 0.2288
##      Balanced Accuracy : 0.6928
##
##      'Positive' Class : Default
##

```

Balanced Decision Tree – Testing




```
## Setting levels: control = NoDefault, case = Default  
## Setting direction: controls < cases
```

Random forest - Imbalanced Training Set

```
##  
## Call:  
##   randomForest(formula = y ~ ., data = train_set)  
##           Type of random forest: classification  
##                       Number of trees: 500  
## No. of variables tried at each split: 4  
##  
##           OOB estimate of  error rate: 18.04%  
## Confusion matrix:  
##             NoDefault Default class.error  
## NoDefault      15492     863  0.05276674  
## Default        2926    1720  0.62978907  
  
oob_accuracy <- 1 - (rf_model$err.rate[nrow(rf_model$err.rate), "OOB"])  
print(paste("OOB Training Accuracy:", oob_accuracy))  
  
## [1] "OOB Training Accuracy: 0.819580019999048"
```

Imbalanced Random Forest Variable Importance



```
# Testing performance
rf_pred_test <- predict(rf_model, test_set)
con_mat <- confusionMatrix(rf_pred_test, test_set$y, mode = 'everything', positive = "Default")
print(con_mat)
```


```
## Confusion Matrix and Statistics
##
##             Reference
## Prediction  NoDefault Default
##   NoDefault      6620     1247
##   Default        389      743
##
##                   Accuracy : 0.8182
##                   95% CI : (0.8101, 0.8261)
##   No Information Rate : 0.7789
##   P-Value [Acc > NIR] : < 2.2e-16
##
##                   Kappa : 0.3759
##
##   Mcnemar's Test P-Value : < 2.2e-16
##
##                   Sensitivity : 0.37337
##                   Specificity  : 0.94450
##   Pos Pred Value  : 0.65636
##   Neg Pred Value : 0.84149
```

```

##                  Precision : 0.65636
##                  Recall   : 0.37337
##                  F1      : 0.47598
##      Prevalence   : 0.22114
##      Detection Rate : 0.08256
## Detection Prevalence : 0.12579
##      Balanced Accuracy : 0.65893
##
##      'Positive' Class : Default
##

```

Imbalanced Random Forest – Testing



```

## Setting levels: control = NoDefault, case = Default
## Setting direction: controls < cases

```

Random forest - Balanced Training Set

```

##
## Call:
##   randomForest(formula = y ~ ., data = train_set_bal)
##   Type of random forest: classification
##   Number of trees: 500
##   No. of variables tried at each split: 4
##

```

```


##          OOB estimate of  error rate: 4.99%
## Confusion matrix:
##             NoDefault Default class.error
## NoDefault      14915    1440  0.08804647
## Default        193     16162  0.01180067

oob_accuracy <- 1 - (rf_model_bal$err.rate[nrow(rf_model_bal$err.rate), "OOB"])
print(paste("OOB Training Accuracy:", oob_accuracy))

## [1] "OOB Training Accuracy: 0.950076429226536"

```

Balanced Random Forest Variable Importance



```

# Balanced Training performance on original training set
rf_pred_train_bal <- predict(rf_model_bal, train_set)
con_mat <- confusionMatrix(rf_pred_train_bal, train_set$y, mode = 'everything', positive = "Default")
print(con_mat)

## Confusion Matrix and Statistics
##
##             Reference
## Prediction  NoDefault Default
##   NoDefault      16195      0
##   Default         160     4646
## 
##           Accuracy : 0.9924

```

```

##                               95% CI : (0.9911, 0.9935)
##      No Information Rate : 0.7788
##      P-Value [Acc > NIR] : < 2.2e-16
##
##                  Kappa : 0.9782
##
##  Mcnemar's Test P-Value : < 2.2e-16
##
##      Sensitivity : 1.0000
##      Specificity : 0.9902
##      Pos Pred Value : 0.9667
##      Neg Pred Value : 1.0000
##                  Precision : 0.9667
##                  Recall : 1.0000
##                  F1 : 0.9831
##      Prevalence : 0.2212
##      Detection Rate : 0.2212
##      Detection Prevalence : 0.2288
##      Balanced Accuracy : 0.9951
##
##      'Positive' Class : Default
##

```

```

# Testing performance
rf_pred_test_bal <- predict(rf_model_bal, test_set)
con_mat <- confusionMatrix(rf_pred_test_bal, test_set$y, mode = 'everything', positive = "Default")
print(con_mat)

```


```

## Confusion Matrix and Statistics
##
##                  Reference
## Prediction  NoDefault Default
##   NoDefault      6379    1082
##   Default        630     908
##
##                  Accuracy : 0.8098
##                               95% CI : (0.8015, 0.8178)
##      No Information Rate : 0.7789
##      P-Value [Acc > NIR] : 3.688e-13
##
##                  Kappa : 0.3988
##
##  Mcnemar's Test P-Value : < 2.2e-16
##
##      Sensitivity : 0.4563
##      Specificity : 0.9101
##      Pos Pred Value : 0.5904
##      Neg Pred Value : 0.8550
##                  Precision : 0.5904
##                  Recall : 0.4563
##                  F1 : 0.5147
##      Prevalence : 0.2211
##      Detection Rate : 0.1009
##      Detection Prevalence : 0.1709

```


```
##      Balanced Accuracy : 0.6832
##      'Positive' Class : Default
##
```

Balanced Random Forest – Testing



```
## Setting levels: control = NoDefault, case = Default
## Setting direction: controls < cases
```

Comparison of ROC Curves – Imbalanced Data



Comparison of ROC Curves – Balanced Data

